

Cross-Residual Graph Neural Networks for Biomolecular Graph Classification: An Empirical Study of Residual Routing

Xuelin Hu¹ and Xiaoqin Fu²

¹School of Railway Communication and Signaling Engineering, Liuzhou Railway Vocational Technical College, Liuzhou 545000, China

²School of Railway Locomotive and Rolling Stock, Liuzhou Railway Vocational Technical College, Liuzhou 545000, China

Corresponding author:

Xiaoqin Fu²

Email address: fuxiaoqin@ltzy.edu.cn

ABSTRACT

Background. While skip connections and hierarchical pooling have been widely used to preserve multi-depth information in graph neural networks, no prior work directly isolates whether cross-branch hidden-state exchange or cross-block graph-summary reinjection is preferable to ordinary same-path residual reuse for biomolecular graph classification. This is the gap addressed here.

Methods. We propose Cross-Residual Graph Neural Networks (CR-GNN), a controlled comparison framework encompassing five information-flow topologies—plain stacking, node-level residual reuse, node-level cross exchange, graph-level residual propagation, and graph-level cross-summary exchange—evaluated under shared backbones and stratified five-fold validation. Targeted ablations on gate control, residual routing, and parameter sweeps, together with depth-scaling analyses of hidden-state similarity and gradient norms ($L = 1-8$), isolate the mechanisms behind the observed performance patterns.

Results. Within the CR-GNN family, graph-level residual propagation achieves the highest peak accuracy most frequently, while node-level cross-residual exchange dominates on large sparse protein graphs under attention-based operators. Cross-residual architectures exhibit the strongest cross-dataset rank consistency and, when combined with sparse routing, form the most frequently optimal configuration. A depth-scaling analysis further reveals a counter-intuitive pattern: the architecture with the weakest representational continuity at depth achieves the most frequent wins, challenging the assumption that stronger layer-wise stability predicts better performance.

Conclusions. CR-GNN is a design space that reveals systematic trade-offs among peak accuracy, cross-dataset consistency, and representational continuity. Rather than a single winning architecture, the contribution is a reproducible framework that makes these tensions quantitatively visible and provides practical guidance for topology and routing selection.

1 INTRODUCTION

Graph neural networks (GNNs) are now a standard approach for learning from molecular and protein graphs because they encode local neighborhood structure together with higher-order relational context (Kipf and Welling, 2016; Gilmer et al., 2017; Hamilton et al., 2017; Wu et al., 2021; Zhou et al., 2020). In biomolecular graph classification, this combination matters because predictive evidence may depend simultaneously on local biochemical motifs and broader graph-level organization. Benchmark suites such as PROTEINS, DD, ENZYMES, MUTAG, AIDS, and Mutagenicity therefore remain useful as controlled testbeds for methodology-oriented graph classification studies rather than as direct proxies for downstream biological discovery.

The main methodological difficulty is that deeper message passing is not automatically beneficial in these settings. Repeated propagation can suppress informative intermediate states through over-smoothing and related depth effects (Li et al., 2018; Oono and Suzuki, 2020; Chen et al., 2020b), while over-squashing can cause information from distant nodes to be compressed into uninformative bottlenecks (Alon and

Yahav, 2021; Topping et al., 2021). Residual connections, normalization, and regularization alleviate some of this degradation (He et al., 2016; Bresson and Laurent, 2017; Zhao and Akoglu, 2020; Cai et al., 2021; Rong et al., 2019; Chen et al., 2020a), but most residual designs still reuse information along the same branch. As a result, they improve optimization without directly testing whether alternative information-flow topologies preserve more useful historical signals.

Related graph-learning work has already explored skip connections, dense aggregation, APPNP-style propagation, Jumping Knowledge readout, and hierarchical pooling (Xu et al., 2018; Klicpera et al., 2018; Gao and Ji, 2019; Ying et al., 2018). These methods preserve information from multiple depths or across pooling hierarchies, yet they do not directly isolate whether exchanging hidden states across branches or reinjecting pooled graph summaries across block boundaries is preferable to ordinary same-path residual reuse. This is the narrower gap addressed in the present paper.

This paper studies Cross-Residual Graph Neural Networks (CR-GNN), a controlled model family that compares plain stacking, node-level residual reuse, node-level cross-residual exchange, graph-level residual propagation, and graph-level cross-summary exchange under shared graph-classification backbones. The study is organized around three specific research questions:

1. **When does cross-residual interaction help?** Under what dataset regimes and task conditions does exchanging hidden states across parallel branches improve graph classification over ordinary same-path reuse?
2. **When does ordinary residual reuse remain the stronger default?** On which biomolecular benchmarks does the simpler single-branch residual design continue to outperform cross-residual alternatives, and what structural properties of those datasets explain this pattern?
3. **How does residual routing design change the answer?** Beyond the binary question of whether an extra reuse path exists, how do different filtering strategies—learnable dense routing, top-k selection, and sparse suppression—affect the relative performance of residual and cross-residual architectures?

Framing the study around these three questions is important for the paper’s methodological scope. Rather than claiming that one additional pathway always yields a better GNN, the goal is to reveal *when and how* residual routing helps—identifying which reuse topology preserves discriminative historical structural signal reliably once backbone, data split, and training protocol are controlled.

The work is organized as a biomolecular graph-classification methods study with a mechanism-oriented empirical focus. The main evidence comes from a protein-oriented benchmark core built around PROTEINS, DD, and ENZYMES, together with supplementary molecular robustness datasets. This design keeps the paper close to biologically meaningful graph settings while maintaining a reproducible and controlled benchmark scope. The later sections therefore move from the architectural motivation to the common methodological setting, then to the confirmatory benchmark results, targeted ablations, and supporting sensitivity analyses. This progression is intentional: the benchmark first establishes the empirical pattern, and the gate and routing studies then test whether that pattern can be explained by how residual information is passed forward.

1.1 Contributions

The main contributions of this work are:

- **A controlled comparison framework that reveals design trade-offs.** Rather than proposing a single new architecture and claiming it is universally better, this work compares five information-flow designs—plain stacking, node-level residual reuse, node-level cross-residual exchange, graph-level residual propagation, and graph-level cross-summary exchange—under identical backbones, data splits, and training protocols. Across 24 dataset-operator combinations, the framework reveals that no single topology dominates all dimensions: GraphResGNN wins 50% of combinations but shows the weakest representational continuity at depth; NodeCrossGNN wins 25% and dominates DD across three of four operators; GraphCrossGNN is the most rank-consistent design. This controlled setup makes it possible to identify *which* reuse strategy to choose under *which* conditions, rather than asking which one wins on average.

- **Empirical evidence for systematic performance trade-offs.** Under stratified five-fold evaluation across six biomolecular datasets and four operators, GraphResGNN achieves the highest peak accuracy most frequently, but its advantage is not universal: it ranks fifth on MUTAG. NodeCrossGNN is the second-strongest architecture, particularly on large sparse protein graphs (DD) where it wins under three of four operators, and on attention-based configurations (GATConv). GraphCrossGNN, while never ranking first, is the most consistent performer across all datasets (rank standard deviation 0.75), consistently placing in the top half. The depth-scaling analysis ($L = 1-8$) further reveals that peak accuracy and representational continuity are not aligned: NodeResGNN achieves the highest CKA at depth (0.968 at $L = 8$) but exhibits the greatest performance volatility across datasets.
- **A depth-scaling mechanism analysis.** Through layer-wise CKA, cosine similarity, and gradient norm tracking from $L = 1$ to $L = 8$ layers, this study provides the mechanism-level evidence behind the performance patterns. NodeResGNN maintains the strongest representational continuity across depth, yet this does not translate into superior peak accuracy. GraphResGNN, despite the lowest CKA at depth, achieves the most frequent wins—suggesting that aggressive representational change between layers, when coupled with graph-level context reinjection, can be a productive rather than destructive design feature. GraphCrossGNN occupies a consistent intermediate position in both performance rank and representational continuity, making it the safest choice when dataset characteristics are uncertain.
- **A routing-level mechanism analysis with practical guidance.** Beyond architecture labels, this study examines *how* residual information is filtered before reuse—comparing learnable dense routing, top-k channel selection, and sparse suppression under the same five-fold protocol. Sparse routing combined with cross-residual architectures emerges as the most frequently optimal configuration, and routing strategy is shown to be target-dependent. This provides practitioners with actionable guidance: the choice of architecture and routing should be matched to dataset characteristics, with cross-residual plus sparse routing recommended for large sparse graphs and attention-based operators, and graph-level residual plus learnable routing for datasets dominated by global context.

2 MATERIALS AND METHODS

This section defines the common graph-classification setting, the CR-GNN architecture family, the benchmark package, and the shared experimental protocol. The guiding principle is to isolate where information is reused while keeping the remaining pipeline as consistent as possible. That design choice is central to the paper’s methodological justification: if backbone operator, pooling, and evaluation protocol are held fixed, then performance differences are more plausibly attributable to the reuse topology itself rather than to unrelated implementation changes.

2.1 Problem definition and notation

Let $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ denote an undirected or directed graph, where $\mathcal{V} = \{v_1, v_2, \dots, v_n\}$ is the set of n nodes and $\mathcal{E} \subseteq \mathcal{V} \times \mathcal{V}$ is the set of edges. Each node $v_i \in \mathcal{V}$ is associated with a feature vector $\mathbf{x}_i \in \mathbb{R}^{d_{\text{in}}}$, and the node features for the entire graph are collected in $\mathbf{X} \in \mathbb{R}^{n \times d_{\text{in}}}$. The graph structure is represented using an adjacency matrix $\mathbf{A}$.

Given a training dataset $\mathcal{D} = \{(\mathcal{G}_i, y_i)\}_{i=1}^N$, the objective is to learn a mapping function $f_\theta : \mathcal{G} \rightarrow \hat{y}$ that predicts the label of an unseen graph. The present study compares five architectures under the same graph-classification pipeline:

- **PlainGNN:** sequential message passing without explicit state reuse
- **NodeResGNN:** single-branch node-level residual reuse
- **NodeCrossGNN:** two-branch node-level cross exchange
- **GraphResGNN:** graph-level residual propagation across sequential blocks
- **GraphCrossGNN:** graph-level cross-summary exchange across paired blocks

The comparison is designed to isolate where information is reused: inside one branch, across parallel branches, or through pooled graph summaries passed between blocks. This isolates a concrete research question behind the architecture family. Same-branch residual reuse mainly tests whether re-injecting historical states stabilizes deeper propagation, whereas cross-residual designs test whether exchanging information across parallel streams can preserve complementary structure that would otherwise be lost under ordinary stacking.

2.2 Shared message passing and graph-level learning

Let Φ denote a message-passing operator. For a chosen backbone, node representations are updated over L layers as

$$\mathbf{h}_i^{(\ell+1)} = \sigma \left(\mathbf{W}^{(\ell)} \cdot \text{AGGREGATE}_{\Phi} \left(\left\{ \mathbf{h}_j^{(\ell)} : j \in \mathcal{N}(i) \right\}, \mathbf{h}_i^{(\ell)} \right) + \mathbf{b}^{(\ell)} \right). \quad (1)$$

After L layers, a readout function R aggregates node-level representations into a graph embedding

$$\mathbf{h}_{\mathcal{G}} = R \left(\left\{ \mathbf{h}_i^{(L)} : i = 1, \dots, n \right\} \right). \quad (2)$$

The classifier produces

$$\hat{\mathbf{y}} = \text{softmax} \left(\mathbf{W}_{\text{clf}} \mathbf{h}_{\mathcal{G}} + \mathbf{b}_{\text{clf}} \right), \quad (3)$$

and the model is trained end-to-end by minimizing cross-entropy:

$$\mathcal{L}(\theta) = -\frac{1}{N} \sum_{i=1}^N \sum_{c=1}^C y_{i,c} \log(\hat{y}_{i,c}). \quad (4)$$

2.3 CR-GNN architecture family

CR-GNN is defined as a compact family of graph classification architectures. All variants share the same input projection, stacked message-passing layers, graph-level pooling, and linear classifier; they differ only in how intermediate node states and graph-level summaries are reused across depth and across branches. The design deliberately avoids introducing separate encoders, task-specific feature engineering, or architecture-specific readout heads, because such additions would make it harder to interpret whether any observed gain comes from residual topology or from extra model capacity.

Let $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ be a graph with node feature matrix $\mathbf{X} \in \mathbb{R}^{n \times d_{\text{in}}}$ and adjacency structure $\mathbf{A}$. The model can be instantiated with one of four standard message-passing operators,

$$\Phi \in \{\text{GCNConv}, \text{GATConv}, \text{SAGEConv}, \text{GINConv}\},$$

and reuses the same operator type throughout a model instance. For a hidden width d , every branch starts with an input projection

$$\mathbf{H}^{(0)} = \text{ReLU}(\Phi_{\text{in}}(\mathbf{X}, \mathbf{A})), \quad (5)$$

followed by L hidden layers and a readout

$$\mathbf{g} = \text{Pool}(\mathbf{H}^{(L)}). \quad (6)$$

2.3.1 Single-branch and node-level residual variants

PlainGNN. The plain baseline stacks L copies of the same operator without residual reuse:

$$\mathbf{H}^{(\ell+1)} = \psi \left(\text{ReLU} \left(\Phi_{\ell} \left(\mathbf{H}^{(\ell)}, \mathbf{A} \right) \right) \right). \quad (7)$$

NodeResGNN. The first residual variant stores the previous hidden state within the same branch:

$$\mathbf{H}^{(\ell+1)} = \psi \left(\text{ReLU} \left(\Phi_{\ell} \left(\mathbf{H}^{(\ell)} + \mathbf{H}^{(\ell-1)}, \mathbf{A} \right) \right) \right). \quad (8)$$

172 **NodeCrossGNN.** The node-level cross-residual model maintains two parallel branches with separate
 173 input projections:

$$\mathbf{H}_1^{(\ell+1)} = \psi\left(\text{ReLU}\left(\Phi_\ell^{(1)}\left(\mathbf{H}_1^{(\ell)} + \tilde{\mathbf{H}}_2^{(\ell)}, \mathbf{A}\right)\right)\right), \quad (9)$$

$$\mathbf{H}_2^{(\ell+1)} = \psi\left(\text{ReLU}\left(\Phi_\ell^{(2)}\left(\mathbf{H}_2^{(\ell)} + \tilde{\mathbf{H}}_1^{(\ell)}, \mathbf{A}\right)\right)\right). \quad (10)$$

174 The final node representation is the branch sum

$$\mathbf{H}^{(L)} = \mathbf{H}_1^{(L)} + \mathbf{H}_2^{(L)}. \quad (11)$$




Figure 1. Node-level CR-GNN architecture family. Left: PlainGNN — sequential message-passing layers without state reuse. Middle: NodeResGNN — single-branch residual reuse, where each layer receives the previous hidden state as additional input. Right: NodeCrossGNN — dual-branch cross-residual exchange, where each branch reuses both its own and the counterpart’s previous hidden state; the two branches are summed before global mean pooling and classification.

175 2.3.2 Graph-level residual propagation

176 **GraphResGNN.** This variant chains three graph-conditioned blocks:

$$\mathbf{g}^{(t+1)} = B_{t+1}\left(\mathbf{X}, \mathbf{A}, \mathbf{g}^{(t)}\right), \quad t \in \{1, 2\}. \quad (12)$$

177 **GraphCrossGNN.** Four blocks are organized as two sequential pairs. Within each stage, the two
 178 block outputs are swapped and reused as the next stage’s graph-level input:

$$\left(\mathbf{g}_1^{(t)}, \mathbf{g}_2^{(t)}\right) = \left(B_{2t-1}\left(\mathbf{X}, \mathbf{A}, \mathbf{g}_1^{(t-1)}\right), B_{2t}\left(\mathbf{X}, \mathbf{A}, \mathbf{g}_2^{(t-1)}\right)\right), \quad (13)$$

$$\left(\mathbf{g}_1^{(t+1)}, \mathbf{g}_2^{(t+1)}\right) = \left(\mathbf{g}_2^{(t)}, \mathbf{g}_1^{(t)}\right). \quad (14)$$

179 The final graph embedding is

$$\mathbf{g}_{\text{final}} = \mathbf{g}_1^{(T)} + \mathbf{g}_2^{(T)}. \quad (15)$$




Figure 2. Graph-level CR-GNN architecture family. Left: GraphResGNN — sequential graph-conditioned blocks with a shared residual state propagating through a residual router across blocks. Right: GraphCrossGNN — dual-branch cross-summary exchange, where two parallel block sequences (Branch A and Branch B) swap graph-level summaries between paired stages. The residual router supports four routing modes: fixed, learnable, top-k, and sparse.

2.4 Residual routing mechanisms

Beyond the architecture family itself, this study evaluates how residual information is injected. This second layer is necessary because the architecture comparison alone cannot answer whether an apparent residual advantage comes from the existence of an extra path or from the way that path is filtered. Three routing modes are used in the targeted ablations:

- **Learnable routing:** dense residual injection modulated by an adaptive gate.
- **Top-k routing:** only the strongest residual channels are retained before injection.
- **Sparse routing:** weak residual coefficients are suppressed, leaving a filtered residual path.

These routing modes are not used to redefine the core architecture family. Instead, they form a mechanism layer on top of selected residual and cross-residual targets in the experimental design. The routing study is therefore interpretive rather than decorative: it tests whether the same architecture behaves differently when residual information is passed forward densely, selectively, or sparsely.

2.5 Datasets

To test the proposed architecture in a biologically meaningful setting, the benchmarks are organized into a unified two-layer package.

The main biological benchmark consists of PROTEINS, DD, and ENZYMES (Fey and Lenssen, 2019; Morris et al., 2020). The supplementary robustness package consists of MUTAG, AIDS, and Mutagenicity. This organization keeps the main biological claim centered on protein-oriented datasets while still testing structural robustness on additional molecular graphs.

PROTEINS PROTEINS contains graph representations of proteins, where nodes denote secondary structure elements and edges encode neighborhood relations between them. The task is binary graph classification. The dataset contains 1,113 graphs, 2 classes, and 3 input features per node.

DD DD is a protein structure dataset in which each graph corresponds to a protein and nodes represent amino acids or structure-related units linked by proximity-derived edges. The task is binary classification. DD contains 1,178 graphs, 2 classes, and 89 input features.

ENZYMES ENZYMES is a 6-class enzyme classification benchmark with 600 graphs and 3 input features. It preserves an enzyme-oriented biological interpretation while remaining directly comparable to the other graph classification benchmarks in the study.

All active datasets are used through a unified benchmark interface (Fey and Lenssen, 2019; Morris et al., 2020). We keep preprocessing intentionally light and do not introduce task-specific feature engineering, graph augmentation, or edge redesign in the final protocol.

All reported results follow the same stratified two-stage evaluation design:

- **Outer evaluation:** stratified 5-fold cross-validation at the graph level
- **Inner validation:** a stratified validation subset sampled from the training fold with validation ratio 0.1

The primary evaluation metric is graph classification accuracy on the held-out test fold. For fold k with test set $\mathcal{T}_k$, the fold accuracy is

$$\text{Acc}_k = \frac{1}{|\mathcal{T}_k|} \sum_{(\mathcal{G}_i, y_i) \in \mathcal{T}_k} \mathbf{1}[\hat{y}_i = y_i]. \quad (16)$$

Across the five folds, we report the mean accuracy

$$\overline{\text{Acc}} = \frac{1}{5} \sum_{k=1}^5 \text{Acc}_k \quad (17)$$

and the corresponding standard deviation

$$\sigma_{\text{Acc}} = \sqrt{\frac{1}{5} \sum_{k=1}^5 (\text{Acc}_k - \overline{\text{Acc}})^2}. \quad (18)$$

Table 1. Unified summary of the main biological benchmark package.

Dataset	Split	#Graphs	#Cls.	Feat.	Avg. N / E	Role
PROTEINS	5-fold CV + val	1113	2	3	39.1 / 72.8	main
DD	5-fold CV + val	1178	2	89	284.3 / 715.7	main
ENZYMES	5-fold CV + val	600	6	3	32.6 / 62.1	main

Table 2. Supplementary robustness datasets retained outside the core biological narrative.

Dataset	Split	#Graphs	#Cls.	Feat.	Avg. N / E	Role
MUTAG	5-fold CV + val	188	2	7	17.9 / 19.8	supp.
AIDS	5-fold CV + val	2000	2	38	15.7 / 16.2	supp.
Mutagenicity	5-fold CV + val	4337	2	14	30.0 / 30.6	supp.

2.6 Compared methods and backbone operators

The full paper-facing benchmark compares nine methods. The CR-GNN family contains five models: PlainGNN, NodeResGNN, NodeCrossGNN, GraphResGNN, and GraphCrossGNN. Four external baselines are also included in the benchmark package: GraphSAGEBaseline, GINBaseline, JKNetBaseline, and APPNPBaseline. These baselines are intended to anchor the controlled comparison against representative graph-classification families with distinct propagation or aggregation behavior, rather than to claim exhaustive coverage of every recent leaderboard method.

Each method is evaluated under four message-passing operators:

- GCNConv
- GATConv
- GINConv
- SAGEConv

2.7 Training configuration and implementation details

All models are implemented in PyTorch Geometric (Fey and Lenssen, 2019) under a unified training protocol. The key hyperparameters are summarized in Table 3. Training is terminated by early stopping when validation loss fails to improve for 60 consecutive epochs, and the checkpoint with the lowest validation loss is retained for test evaluation. No batch normalization or layer normalization is applied; the only regularization is dropout and weight decay. All train-validation-test splits are fixed by a global random seed of 1024, ensuring identical data partitions across models within each fold. No task-specific feature engineering, graph augmentation, or edge redesign is applied; the raw node features and adjacency structures from TUDataset are used directly.

The learnable gate used in residual and cross-residual variants applies a sigmoid activation to a scalar parameter initialized at 0.8, allowing each residual injection to be adaptively scaled during training. For the top-k routing mode, only the strongest fraction of residual channels (determined by absolute magnitude) is retained before injection. For the sparse routing mode, a softshrink operator with threshold λ suppresses residual coefficients whose magnitude falls below λ , leaving a filtered residual path.

2.8 Experimental design

The empirical study is organized to separate confirmatory five-fold evidence from exploratory supporting analyses. The latest benchmark evaluates six datasets, nine methods, four message-passing operators, and five outer folds under one unified protocol. This full matrix provides the main comparison used throughout the paper, and the benchmark tables are placed directly in the corresponding result subsections so that each empirical claim is supported in place.

The benchmark is divided into two layers. The biological core consists of PROTEINS, DD, and ENZYMES, which anchor the main biomolecular interpretation of the paper. A supplementary robustness package extends the analysis to MUTAG, AIDS, and Mutagenicity. This split keeps the central claim focused while still testing whether the conclusions remain stable on additional molecular graph settings.

Table 3. Key training hyperparameters shared across all experiments.

Parameter	Value
Framework	PyTorch Geometric
Hidden dimension (d)	64
Message-passing layers (L)	4
Dropout rate	0.5
Optimizer	Adam
Initial learning rate	5×10^{-3}
Weight decay	1×10^{-4}
LR schedule	ReduceLROnPlateau (factor 0.5, patience 15)
Minimum learning rate	1×10^{-5}
Batch size	256
Maximum epochs	200
Early stopping patience	60 epochs
Gradient clip norm (L_2)	2.0
Pooling	Global mean pooling
Loss function	Cross-entropy
Random seed	1024
Validation split ratio	0.1

Targeted ablations are organized to address the three research questions posed in the Introduction through controlled mechanism experiments:

- Q1: How do reuse topologies trade off peak performance, consistency, and representational stability?** *Full-benchmark and depth-scaling analyses.* All five architectures are compared across six datasets and four operators under stratified five-fold evaluation, and hidden-state similarity (CKA/cosine) together with per-layer gradient norms are tracked from $L = 1$ to $L = 8$ layers on PROTEINS under GCNConv. This isolates whether stronger raw performance comes at the cost of representational continuity, or whether certain topologies can achieve both.
- Q2: When does cross-residual interaction outperform residual reuse?** *Gate-control and cross-gate ablations.* On representative residual-oriented and cross-oriented targets, learnable versus fixed-strength gates are compared under identical five-fold conditions. This directly tests whether the cross-residual advantage on specific dataset–operator combinations (notably DD with GATConv/SAGEConv/GINConv) depends on adaptive injection control, or whether any fixed-weight additional path produces similar gains.
- Q3: How does residual routing design change the answer?** *Residual-routing ablations.* Four representative targets (spanning both residual-oriented and cross-oriented settings) compare learnable dense routing, top-k channel selection at three retention ratios (0.25, 0.5, 0.75), and sparse suppression at three thresholds (0.02, 0.05, 0.10) under the same five-fold protocol. This moves the analysis beyond architecture labels—whether a model is called “residual” or “cross-residual”—to the mechanism level: how aggressively historical information should be filtered before reuse.

Residual-parameter sweeps extend the routing analysis by varying top-k ratios and sparsity strengths on the same representative targets. Older single-fold sensitivity scans over depth, dropout, and learning rate are retained only as exploratory supporting analyses. In the final interpretation, the five-fold latest-version runs are treated as the confirmatory evidence, while the single-fold scans are used only to describe optimization trends rather than to support the main claim.

2.9 Evaluation metrics and statistical reporting

The primary metric is graph classification accuracy on the held-out test fold. Main benchmark and targeted ablation results are reported as five-fold mean accuracy with standard deviation. In the final paper, the five-fold latest-version experiments are treated as the main confirmatory evidence, while older single-fold parameter scans are used only as supporting exploratory analyses.

3 RESULTS

This section reports empirical evidence in four layers: the full benchmark, depth-scaling mechanism analyses, internal CR-GNN trade-off quantification, and targeted ablation studies. The main benchmark and ablations use stratified five-fold results reported as mean best test accuracy $\pm$ fold-level standard deviation; older single-fold sensitivity scans are retained only as supporting analyses.

3.1 Overall benchmark performance

The final benchmark evaluates six datasets and nine methods (five CR-GNN variants plus four external baselines), with CR-GNN models using GCNConv and each baseline using its native operator, under five outer folds.

The three tables together reveal a nuanced landscape. Within the CR-GNN family, GraphResGNN is the strongest on PROTEINS, ENZYMES, and AIDS; NodeResGNN leads on PROTEINS. Among external baselines, GIN is particularly strong on MUTAG (0.8403) and AIDS (0.9145), while GraphSAGE leads on DD (0.7334) and Mutagenicity (0.8084). The presence of competitive external baselines reinforces the paper’s central framing: the CR-GNN family is not claimed to be universally dominant, but rather provides a controlled design space for studying systematic trade-offs.

3.2 Results on the biological core datasets

The main biological benchmark results are reported in Table 4.

PROTEINS. Within the CR-GNN family, NodeResGNN leads numerically (0.7143 ± 0.0223), followed closely by GraphResGNN (0.7071 ± 0.0402). External baselines GraphSAGE (0.6954), GIN (0.6972), and JKNet (0.6954) cluster near the CR-GNN models. APPNP performs markedly worse (0.5966).

DD. GraphSAGE leads overall (0.7334 ± 0.0227), with GraphResGNN as the strongest CR-GNN variant (0.7275 ± 0.0144). Notably, the two cross-residual variants—GraphCrossGNN (0.7207) and NodeCrossGNN (0.7164)—both outperform NodeResGNN (0.7088), supporting the observation that cross-branch exchange is especially effective on large, sparse protein graphs.

ENZYMES. All models perform poorly in absolute terms, with GraphResGNN leading (0.3333 ± 0.1000). The CR-GNN family’s relative ordering (GraphRes $>$ GraphCross $>$ NodeCross $\approx$ NodeRes $\approx$ Plain) is consistent with patterns on other datasets, though the 6-class, 600-graph setting creates a noisy regime that limits the strength of architecture-level conclusions.

3.3 Results on the supplementary robustness datasets

The supplementary robustness results are reported in Table 5.

MUTAG. GIN is the strongest model overall (0.8403 ± 0.0298), substantially ahead of all other methods. Within the CR-GNN family, NodeResGNN (0.7451) and NodeCrossGNN (0.7398) are competitive with GraphSAGE (0.7504). GraphResGNN performs worst (0.7238), consistent with the observation that aggressive representational reshaping is less suited to small molecular graphs.

AIDS. GIN (0.9145 ± 0.0415) and GraphResGNN (0.9090 ± 0.0194) are nearly tied, with GraphResGNN showing substantially lower variance. GraphSAGE (0.8800), NodeCrossGNN (0.8850), JKNet (0.8775), and GraphCrossGNN (0.8740) form a tight cluster.

Mutagenicity. GraphSAGE leads (0.8084 ± 0.0180), followed closely by GraphResGNN (0.8022 ± 0.0244) and GraphCrossGNN (0.7992 ± 0.0232). The top four models are separated by only 0.015, with overlapping standard deviations.

3.4 Depth scaling and representation similarity

To understand the mechanism behind the benchmark patterns, we track hidden-state similarity (CKA and cosine) and per-layer gradient norms across $L = 1$ through $L = 8$ layers on PROTEINS under GCNConv.

Table 4. Main biological benchmark results. Cells show mean $\pm$ std [min, max] over five folds. Bold: best mean per dataset.

Model	PROTEINS	DD	ENZYMES
Plain	0.7026 $\pm$ 0.0331 [0.655, 0.735]	0.6970 $\pm$ 0.0421 [0.646, 0.763]	0.2683 $\pm$ 0.0509 [0.217, 0.358]
NodeRes	0.7143 $\pm$ 0.0223 [0.685, 0.744]	0.7088 $\pm$ 0.0332 [0.651, 0.754]	0.2683 $\pm$ 0.0517 [0.217, 0.367]
NodeCross	0.6945 $\pm$ 0.0488 [0.613, 0.735]	0.7164 $\pm$ 0.0315 [0.677, 0.771]	0.2717 $\pm$ 0.0692 [0.167, 0.367]
GraphRes	0.7071 $\pm$ 0.0402 [0.655, 0.749]	0.7275 $\pm$ 0.0144 [0.711, 0.747]	0.3333 $\pm$ 0.1000 [0.150, 0.425]
GraphCross	0.7017 $\pm$ 0.0520 [0.608, 0.757]	0.7207 $\pm$ 0.0322 [0.694, 0.784]	0.2850 $\pm$ 0.0868 [0.167, 0.408]
GraphSAGE	0.6954 $\pm$ 0.0253 [0.659, 0.731]	0.7334 $\pm$ 0.0227 [0.694, 0.763]	0.2717 $\pm$ 0.0662 [0.192, 0.358]
GIN	0.6972 $\pm$ 0.0364 [0.635, 0.740]	0.7105 $\pm$ 0.0160 [0.689, 0.733]	0.2883 $\pm$ 0.1148 [0.158, 0.492]
JKNet	0.6954 $\pm$ 0.0514 [0.599, 0.739]	0.7122 $\pm$ 0.0304 [0.677, 0.758]	0.2667 $\pm$ 0.0660 [0.158, 0.350]
APPNP	0.5966 $\pm$ 0.0014 [0.595, 0.599]	0.7003 $\pm$ 0.0220 [0.668, 0.726]	0.2133 $\pm$ 0.0239 [0.175, 0.242]

Table 5. Supplementary robustness benchmark results. Cells show mean $\pm$ std [min, max] over five folds. Bold: best mean per dataset.

Model	MUTAG	AIDS	Mutagenicity
Plain	0.7343 $\pm$ 0.0433 [0.684, 0.811]	0.8295 $\pm$ 0.0456 [0.755, 0.877]	0.7830 $\pm$ 0.0170 [0.764, 0.812]
NodeRes	0.7451 $\pm$ 0.0504 [0.684, 0.838]	0.8615 $\pm$ 0.0314 [0.800, 0.885]	0.7874 $\pm$ 0.0142 [0.769, 0.809]
NodeCross	0.7398 $\pm$ 0.0523 [0.684, 0.838]	0.8850 $\pm$ 0.0380 [0.818, 0.930]	0.7939 $\pm$ 0.0196 [0.770, 0.826]
GraphRes	0.7238 $\pm$ 0.0778 [0.649, 0.865]	0.9090 $\pm$ 0.0194 [0.890, 0.935]	0.8022 $\pm$ 0.0244 [0.763, 0.832]
GraphCross	0.7397 $\pm$ 0.0531 [0.684, 0.838]	0.8740 $\pm$ 0.0460 [0.800, 0.943]	0.7992 $\pm$ 0.0232 [0.756, 0.826]
GraphSAGE	0.7504 $\pm$ 0.0506 [0.684, 0.838]	0.8800 $\pm$ 0.0537 [0.780, 0.930]	0.8084 $\pm$ 0.0180 [0.786, 0.827]
GIN	0.8403 $\pm$ 0.0298 [0.811, 0.895]	0.9145 $\pm$ 0.0415 [0.838, 0.960]	0.7985 $\pm$ 0.0189 [0.777, 0.822]
JKNet	0.7292 $\pm$ 0.0600 [0.658, 0.838]	0.8775 $\pm$ 0.0416 [0.800, 0.925]	0.7897 $\pm$ 0.0243 [0.755, 0.822]
APPNP	0.7131 $\pm$ 0.0603 [0.632, 0.811]	0.8000 $\pm$ 0.0000 [0.800, 0.800]	0.6689 $\pm$ 0.0232 [0.653, 0.714]




Figure 3. CKA similarity between adjacent depth stages across layer counts $L = 1$ through $L = 8$. Each subplot shows one model; curves correspond to different layer counts. Higher CKA indicates greater representational continuity across depth.

Table 6. Winner summary with max/min fold accuracy and parameter count for each dataset.

Dataset	Winner	Mean Acc	Max Acc	Min Acc	Best Epoch	Params
PROTEINS	NodeRes	0.7143	0.7444	0.6847	69.0	17027
DD	GraphSAGE	0.7334	0.7627	0.6936	55.8	44610
ENZYMES	GraphRes	0.3333	0.4250	0.1500	70.6	52251
MUTAG	GIN	0.8403	0.8947	0.8108	134.2	38082
AIDS	GIN	0.9145	0.9600	0.8375	102.8	40066
Mutagenicity	GraphSAGE	0.8084	0.8270	0.7860	122.2	35010




Figure 4. Cosine similarity between adjacent depth stages across layer counts $L = 1$ through $L = 8$. Negative values indicate anti-correlated representations between adjacent stages.

Figure 3 shows the CKA similarity between adjacent depth stages for each model, and Figure 4 shows the corresponding cosine similarity. Three patterns emerge. First, deeper stacking consistently reduces adjacent-stage similarity across all five architectures, confirming that depth-induced representation divergence is a general phenomenon. Second, at $L = 2\text{--}3$, all architectures maintain relatively high CKA; at $L = 8$, the ordering becomes NodeResGNN (0.968) > GraphCrossGNN (0.952) > PlainGNN (0.951) > NodeCrossGNN (0.946) > GraphResGNN (0.942). Third, and most importantly, the relationship between CKA and performance is the opposite of what a naive representational-stability hypothesis would predict: the model with the highest CKA (NodeResGNN) wins only 3 of 24 combinations, while the model with the lowest CKA (GraphResGNN) wins 12 of 24. This counter-intuitive pattern suggests that GraphResGNN’s sequential block structure, by deliberately reshaping representations at each stage using reinjected graph-level context, produces more discriminative final representations precisely because of—rather than in spite of—the representational change between stages.




Figure 5. Deepest-layer gradient L2 norm as a function of layer count ($L = 1$ through $L = 8$) for all five architectures on PROTEINS (GCNConv).

Figure 5 shows that PlainGNN’s deepest gradient norm degrades most rapidly with depth, while the residual family maintains more stable gradients. NodeResGNN and NodeCrossGNN maintain comparable gradient stability across all depths, confirming that cross-branch exchange does not introduce destructive gradient interference. GraphResGNN and GraphCrossGNN occupy intermediate positions.

3.5 Performance trade-offs within the CR-GNN family

To isolate the systematic trade-offs among the five CR-GNN topologies, Table 7 summarizes winner counts across all 24 dataset–operator combinations (GCNConv-only). GraphResGNN achieves the highest peak accuracy in 12 of 24 combinations (50%). NodeCrossGNN is second with 6 wins (25%), dominating DD across three of four operators (GATConv, SAGEConv, GINConv). NodeResGNN wins 3 combinations, PlainGNN wins 2 (both on MUTAG), and GraphCrossGNN wins 1.

Table 7. Winner counts across all 24 dataset–operator combinations (GCNConv-only, CR-GNN family).

Model	PROTEINS	DD	ENZYMES	MUTAG	AIDS	Mutagenicity	Total
GraphResGNN	3	1	3	0	3	2	12
NodeCrossGNN	0	3	1	0	1	1	6
NodeResGNN	1	0	0	1	0	1	3
PlainGNN	0	0	0	2	0	0	2
GraphCrossGNN	0	0	0	1	0	0	1

Beyond win counts, the architectures differ systematically in cross-dataset consistency (Table 8). GraphCrossGNN, despite winning only one combination, is the most rank-stable design (rank std 0.75), never placing below fourth on any dataset. GraphResGNN, while the most frequent winner, drops to fifth on MUTAG (rank std 1.46). NodeCrossGNN occupies an intermediate consistency position (rank std 1.00) and is the closest competitor to the best model on average (mean gap 0.022).

Table 8. Mean rank, average gap to best, and rank stability across six datasets (CR-GNN family, GCNConv).

Model	Mean rank	Mean gap to best	Rank std
GraphResGNN	1.8	0.005	1.46
GraphCrossGNN	2.7	0.018	0.75
NodeCrossGNN	3.0	0.022	1.00
NodeResGNN	3.2	0.024	1.57
PlainGNN	4.3	0.036	0.75

Taken together, the depth-scaling CKA results and the rank-consistency analysis define a three-way tension: GraphResGNN aggressively reshapes representations to achieve the highest peak accuracy but with inconsistent ranking; NodeResGNN preserves representations most faithfully (highest CKA) but with volatile performance; and GraphCrossGNN balances these forces to achieve the most consistent, though not peak, performance. The DD case further illustrates this trade-off: NodeCrossGNN wins under three of four operators, suggesting that dual-branch node-level exchange is especially effective on large, sparse protein graphs where complementary parallel pathways preserve discriminative signals that a single branch would suppress. By contrast, on smaller, denser graphs such as AIDS, GraphResGNN’s sequential graph-conditioned block structure is more effective.

3.6 Ablation evidence: gates, routing, and parameters

3.6.1 Gate ablation results

The gate ablations test whether adaptive control of residual injection is necessary for the observed architectural advantages. On the strongest graph-level residual target (AIDS + GraphResGNN + GINConv), the learnable gate is the strongest setting, though fixed coefficient 0.5 remains close. On the six cross-favored targets, learnable gating is best on three and fixed gates are best on the other three—indicating that cross-residual architectures are not overly sensitive to gate configuration, consistent with their greater rank stability.

Tables 12 and 13 provide the complete paired-test matrix. Most architecture-level numerical gaps are modest relative to fold-level variation, confirming that differences between reuse topologies are real but conditional.

3.6.2 Residual-routing results

The residual-routing analysis adds a mechanism dimension to the trade-off framework. Four representative targets compare learnable, top-k, and sparse routing modes. Sparse routing with moderate threshold (0.05) is the most frequently optimal configuration, winning on three of four targets. Learnable dense routing remains best only on the strongest graph-level residual line. Two of the three sparse-favoring targets involve cross-residual architectures, suggesting that sparse filtering and cross-branch exchange form a particularly effective combination: sparsity suppresses weak residual channels, while the cross-branch pathway ensures that genuinely complementary information is preserved.

3.6.3 Residual-parameter sweep results

The additional routing settings refine rather than change the pattern. The strongest top-k ratio varies by target, confirming no globally optimal retention level. Aggressive sparsity (0.10) damages the strongest graph-level residual target, reinforcing that the benefit of sparsity is conditional on an appropriate threshold.

3.7 Exploratory sensitivity analysis

Auxiliary sensitivity checks over depth, dropout, and learning rate on PROTEINS, DD, and ENZYMES do not materially change the main conclusions. These checks are reported for transparency but should not be read as evidence at the same level as the repeated-split benchmark and targeted ablations.

4 DISCUSSION

The CR-GNN design space reveals that peak accuracy, cross-dataset consistency, and representational continuity are not aligned—and that no single topology dominates all three dimensions. This section draws out the practical implications of these trade-offs.

4.1 The accuracy–consistency–continuity trade-off

The three-way tension documented in Section 3 has a clear structure. GraphResGNN achieves the most frequent wins (50% of combinations) through aggressive representational reshaping: each of its three graph-conditioned blocks reinjects the accumulated graph-level context, deliberately altering the representation at every stage. This produces the lowest CKA at depth (0.942 at $L = 8$) but the highest peak accuracy. The trade-off is brittleness: GraphResGNN drops to fifth on MUTAG, where local motif preservation dominates over global context.

NodeResGNN occupies the opposite corner of the design space, with the highest CKA (0.968) but the greatest performance volatility (rank std 1.57). Preserving representations too faithfully may limit the

Table 9. AIDS gate ablation on GraphResGNN + GINConv. Numbers are five-fold mean accuracy $\pm$ standard deviation. Bold: best per column.

Gate setting	Mean accuracy
Learnable	0.9620 $\pm$ 0.0089
Fixed 0.0	0.9150 $\pm$ 0.0138
Fixed 0.5	0.9605 $\pm$ 0.0144
Fixed 1.0	0.9480 $\pm$ 0.0300

Table 10. Fold-level AIDS gate ablation on AIDS + GraphResGNN + GINConv. Bold: best per column.

Gate setting	Fold 0	Fold 1	Fold 2	Fold 3	Fold 4	Mean $\pm$ Std
Learnable	0.9500	0.9700	0.9700	0.9675	0.9525	0.9620 $\pm$ 0.0089
Fixed 0.0	0.9225	0.9025	0.9125	0.9000	0.9375	0.9150 $\pm$ 0.0138
Fixed 0.5	0.9325	0.9625	0.9675	0.9675	0.9725	0.9605 $\pm$ 0.0144
Fixed 1.0	0.8925	0.9450	0.9700	0.9775	0.9550	0.9480 $\pm$ 0.0300

model’s capacity to adapt to task-specific structure at each depth stage—a “CKA–performance paradox” in which stronger representational continuity does not predict better classification.

Cross-residual architectures occupy a productive middle ground. GraphCrossGNN achieves the second-highest CKA at depth (0.952) while being simultaneously the most rank-consistent performer (rank std 0.75). Its paired-block structure with cross-summary exchange provides enough continuity to avoid the brittleness of GraphResGNN, and enough representational change to remain competitive. The practical implication is that GraphCrossGNN is the safest choice when dataset characteristics are uncertain or when stable performance across heterogeneous tasks is valued over peak single-dataset accuracy.

4.2 When cross-residual interaction dominates

NodeCrossGNN’s performance on DD—winning under three of four operators—is the clearest case of cross-residual dominance in the benchmark. DD is the largest and sparsest protein graph in the biological core (1,178 graphs, 89-dim features), a regime where a single propagation branch is likely to suppress rare but discriminative structural signals under repeated message passing. Dual-branch node-level exchange provides complementary parallel pathways that preserve these signals. The same pattern is visible, though less pronounced, on other GATConv-based configurations (AIDS, Mutagenicity, ENZYMES), where NodeCrossGNN is consistently competitive.

This dataset–operator specificity suggests a practical decision rule: cross-residual architectures, particularly with attention-based operators, are the preferred choice for large, sparse graphs where complementary pathway preservation matters; graph-level residual architectures are preferred for smaller, denser graphs where global context is strongly discriminative.

4.3 Routing as a trade-off modulator

The routing results add an actionable dimension to the architecture-level trade-off framework. Sparse routing with moderate threshold (0.05) is optimal on three of four representative targets, and two of these three involve cross-residual architectures. The mechanism-level synergy is that sparsity suppresses weak residual channels that would otherwise accumulate as noise across depth, while the cross-branch pathway preserves genuinely complementary information through an alternative route. Learnable dense routing remains best on the strongest graph-level residual target (AIDS + GraphResGNN + GINConv), where rich historical context benefits from unfiltered reuse.

The practical guidance that emerges is: combine sparse routing with cross-residual architectures for balanced, consistent performance; use learnable dense routing with graph-level residual architectures when maximizing peak accuracy on well-characterized datasets.

4.4 Why ENZYMES remains difficult

A notable negative result is that all models perform poorly on ENZYMES in absolute terms: the best model (GraphResGNN, 0.333) is only $\sim$ 16 percentage points above the 6-class random baseline (16.7%). Several factors contribute: only 600 graphs with 3-dimensional input features; a 6-class structure creating

Table 11. Cross-gate ablation on six cross-favored targets. Numbers are five-fold mean accuracy $\pm$ standard deviation. Bold: best per row.

Dataset	Target	Learnable	Fixed 0.0	Fixed 0.5	Fixed 1.0
AIDS	NodeCross + GATConv	0.9045 $\pm$ 0.0165	0.9075 $\pm$ 0.0065	0.9195 $\pm$ 0.0114	0.9105 $\pm$ 0.0241
DD	NodeCross + GATConv	0.7173 $\pm$ 0.0271	0.7053 $\pm$ 0.0361	0.7156 $\pm$ 0.0225	0.7147 $\pm$ 0.0288
DD	NodeCross + GINConv	0.7071 $\pm$ 0.0306	0.7012 $\pm$ 0.0198	0.7232 $\pm$ 0.0207	0.7113 $\pm$ 0.0277
ENZYMES	NodeCross + GATConv	0.2800 $\pm$ 0.0629	0.2850 $\pm$ 0.0627	0.2550 $\pm$ 0.0674	0.2650 $\pm$ 0.0690
MUTAG	GraphCross + GATConv	0.7558 $\pm$ 0.0495	0.7451 $\pm$ 0.0407	0.7450 $\pm$ 0.0416	0.7504 $\pm$ 0.0779
Mutagenicity	NodeCross + GATConv	0.8114 $\pm$ 0.0110	0.8033 $\pm$ 0.0192	0.8010 $\pm$ 0.0130	0.8036 $\pm$ 0.0078

Table 12. Paired t -tests for main biological benchmark datasets. Cells: Δ (Cohen’s d , p). **Bold:** $p < 0.05$ ($n=5$ folds).

Dataset	Comparison	GCNConv	GATConv	SAGEConv	GINConv
PROTEINS	NodeRes – Plain	+0.012 (+0.55, 0.286)	+0.000 (+0.00, 0.999)	+0.006 (+0.31, 0.524)	+0.012 (+0.56, 0.277)
	NodeCross – Plain	-0.008 (-0.28, 0.571)	-0.011 (-0.22, 0.647)	-0.002 (-0.07, 0.881)	-0.008 (-0.23, 0.632)
	GraphRes – Plain	+0.004 (+0.45, 0.375)	+0.015 (+0.48, 0.345)	+0.030 (+1.69, 0.019)	+0.040 (+1.10, 0.069)
	GraphCross – Plain	-0.001 (-0.02, 0.961)	-0.008 (-0.24, 0.623)	+0.004 (+0.24, 0.622)	+0.004 (+0.16, 0.743)
	NodeCross – NodeRes	-0.020 (-0.53, 0.302)	-0.011 (-0.19, 0.690)	-0.008 (-0.72, 0.181)	-0.020 (-0.87, 0.125)
	GraphCross – GraphRes	-0.005 (-0.15, 0.751)	-0.023 (-2.18, 0.008)	-0.026 (-1.47, 0.031)	-0.036 (-2.36, 0.006)
	GraphRes – NodeRes	-0.007 (-0.26, 0.588)	+0.015 (+0.37, 0.459)	+0.023 (+0.75, 0.171)	+0.028 (+1.05, 0.079)
	NodeCross – GraphCross	-0.007 (-0.49, 0.337)	-0.003 (-0.10, 0.838)	-0.005 (-0.21, 0.659)	-0.012 (-0.39, 0.427)
DD	NodeRes – Plain	+0.012 (+0.39, 0.429)	+0.002 (+0.24, 0.624)	+0.016 (+0.48, 0.347)	+0.001 (+0.03, 0.946)
	NodeCross – Plain	+0.019 (+0.55, 0.286)	+0.035 (+0.90, 0.115)	+0.023 (+0.87, 0.124)	+0.025 (+0.93, 0.107)
	GraphRes – Plain	+0.030 (+0.64, 0.223)	+0.020 (+0.50, 0.326)	+0.021 (+0.91, 0.112)	+0.002 (+0.06, 0.892)
	GraphCross – Plain	+0.024 (+0.87, 0.123)	+0.019 (+0.46, 0.357)	+0.004 (+0.19, 0.699)	+0.008 (+0.20, 0.672)
	NodeCross – NodeRes	+0.008 (+0.43, 0.395)	+0.033 (+0.88, 0.121)	+0.007 (+0.64, 0.226)	+0.024 (+0.94, 0.103)
	GraphCross – GraphRes	-0.007 (-0.23, 0.637)	-0.002 (-0.11, 0.815)	-0.017 (-2.82, 0.003)	+0.006 (+0.14, 0.762)
	GraphRes – NodeRes	+0.019 (+0.64, 0.228)	+0.019 (+0.46, 0.359)	+0.005 (+0.24, 0.624)	+0.001 (+0.03, 0.954)
	NodeCross – GraphCross	-0.004 (-0.31, 0.526)	+0.016 (+1.07, 0.075)	+0.019 (+1.53, 0.027)	+0.017 (+0.57, 0.274)
ENZYMES	NodeRes – Plain	-0.000 (-0.00, 1.000)	+0.005 (+0.19, 0.697)	-0.017 (-0.94, 0.103)	+0.030 (+0.47, 0.351)
	NodeCross – Plain	+0.003 (+0.09, 0.847)	+0.032 (+0.45, 0.369)	+0.007 (+0.24, 0.614)	+0.032 (+2.12, 0.009)
	GraphRes – Plain	+0.020 (+0.92, 0.107)	+0.020 (+0.32, 0.507)	+0.020 (+0.47, 0.360)	+0.068 (+2.45, 0.002)
	GraphCross – Plain	+0.020 (+0.55, 0.280)	+0.007 (+0.08, 0.875)	+0.020 (+0.33, 0.502)	+0.028 (+0.52, 0.303)
	NodeCross – NodeRes	+0.003 (+0.08, 0.878)	+0.027 (+0.36, 0.471)	+0.023 (+0.42, 0.399)	+0.002 (+0.07, 0.884)
	GraphCross – GraphRes	-0.000 (-0.01, 0.981)	-0.013 (-0.21, 0.672)	+0.000 (+0.00, 1.000)	-0.040 (-0.62, 0.237)
	GraphRes – NodeRes	+0.020 (+0.57, 0.264)	+0.015 (+0.28, 0.562)	+0.037 (+0.82, 0.151)	+0.038 (+1.02, 0.079)
	NodeCross – GraphCross	-0.017 (-0.47, 0.355)	+0.025 (+0.68, 0.211)	-0.013 (-0.40, 0.444)	+0.003 (+0.07, 0.894)

a harder decision boundary than the binary tasks dominating the rest of the benchmark; and the largest fold-level standard deviation (± 0.100) in the biological core. The ENZYMES results are best interpreted as evidence that the architecture ranking remains consistent even under noisy, low-data regimes, rather than as evidence that any architecture solves the task well.

4.5 Evidence hierarchy and study limitations

The strongest claims are supported by the five-fold benchmark, the five-fold targeted ablations, and the $L = 1-8$ depth-scaling analysis. Older single-fold sensitivity scans are only supportive. A second limitation is statistical power: with five paired observations per combination, the tests have limited power to detect small-to-medium effects (Cohen’s $d < 0.5$). A third limitation is baseline scope: the benchmark includes representative baselines but is not an exhaustive leaderboard. A fourth limitation is dataset scale: the study supports conclusions about controlled graph-classification behavior and design trade-offs, not about large-scale scalability or direct biological discovery. A fifth limitation specific to the depth-scaling analysis is that it was conducted on a single dataset-operator combination (PROTEINS, GCNConv); extending it to other settings would test the generality of the CKA-performance paradox.

4.6 Future directions

The most important extension is to test whether the accuracy-consistency-continuity trade-off generalizes to larger graph benchmarks and richer biological inputs. The CKA-performance paradox warrants further investigation across diverse domains to determine whether it is a general property of deep GNN design spaces or specific to the biomolecular setting. The finding that cross-residual plus sparse routing forms a particularly effective combination should be tested on broader tasks where graph topology must be combined with stronger domain supervision.

Table 13. Paired t -tests for supplementary robustness datasets. Cells: Δ (Cohen’s d , p). **Bold:** $p < 0.05$ ($n=5$ folds).

Dataset	Comparison	GCNConv	GATConv	SAGEConv	GINConv
MUTAG	NodeRes – Plain	+0.011 (+0.29, 0.545)	+0.000 (+0.00, 1.000)	-0.016 (-0.40, 0.416)	-0.005 (-0.12, 0.812)
	NodeCross – Plain	+0.016 (+0.36, 0.477)	+0.005 (+0.09, 0.850)	+0.005 (+0.15, 0.758)	+0.016 (+0.34, 0.498)
	GraphRes – Plain	+0.005 (+0.09, 0.864)	+0.011 (+0.33, 0.513)	-0.011 (-0.10, 0.839)	-0.016 (-0.22, 0.670)
	GraphCross – Plain	-0.011 (-0.20, 0.683)	+0.027 (+0.59, 0.262)	-0.022 (-0.28, 0.576)	-0.005 (-0.16, 0.756)
	NodeCross – NodeRes	+0.005 (+0.10, 0.836)	+0.005 (+0.10, 0.845)	+0.022 (+0.33, 0.516)	+0.022 (+0.37, 0.464)
	GraphCross – GraphRes	-0.016 (-0.34, 0.515)	+0.016 (+0.37, 0.467)	-0.011 (-0.26, 0.613)	+0.011 (+0.25, 0.618)
	GraphRes – NodeRes	-0.005 (-0.12, 0.808)	+0.011 (+0.17, 0.742)	+0.005 (+0.08, 0.881)	-0.011 (-0.12, 0.824)
	NodeCross – GraphCross	+0.027 (+0.71, 0.186)	-0.022 (-0.61, 0.240)	+0.027 (+0.53, 0.310)	+0.022 (+0.48, 0.349)
AIDS	NodeRes – Plain	+0.023 (+1.19, 0.062)	+0.010 (+0.89, 0.112)	+0.018 (+0.86, 0.125)	+0.015 (+0.59, 0.261)
	NodeCross – Plain	+0.028 (+1.06, 0.083)	+0.018 (+0.80, 0.144)	+0.020 (+1.52, 0.022)	+0.018 (+0.67, 0.222)
	GraphRes – Plain	+0.070 (+2.41, 0.005)	+0.025 (+1.57, 0.041)	+0.028 (+2.45, 0.004)	+0.048 (+2.47, 0.004)
	GraphCross – Plain	+0.025 (+0.94, 0.089)	+0.013 (+1.10, 0.068)	+0.018 (+1.16, 0.067)	+0.023 (+0.86, 0.124)
	NodeCross – NodeRes	+0.005 (+0.41, 0.420)	+0.008 (+0.52, 0.311)	+0.003 (+0.14, 0.778)	+0.003 (+1.61, 0.005)
	GraphCross – GraphRes	-0.045 (-1.46, 0.043)	-0.013 (-0.66, 0.218)	-0.010 (-0.44, 0.398)	-0.025 (-0.87, 0.125)
	GraphRes – NodeRes	+0.047 (+2.62, 0.002)	+0.015 (+1.19, 0.063)	+0.010 (+0.46, 0.370)	+0.033 (+2.02, 0.014)
	NodeCross – GraphCross	-0.003 (-0.10, 0.848)	+0.005 (+0.44, 0.387)	+0.003 (+0.20, 0.682)	-0.005 (-0.16, 0.751)
Mutagenicity	NodeRes – Plain	+0.004 (+0.26, 0.597)	+0.011 (+0.95, 0.101)	+0.014 (+0.93, 0.106)	+0.013 (+0.83, 0.137)
	NodeCross – Plain	+0.011 (+0.68, 0.201)	+0.020 (+1.05, 0.079)	+0.013 (+0.86, 0.127)	+0.017 (+0.95, 0.101)
	GraphRes – Plain	+0.019 (+1.24, 0.050)	+0.008 (+0.57, 0.273)	+0.012 (+0.62, 0.237)	+0.018 (+0.83, 0.139)
	GraphCross – Plain	+0.016 (+0.85, 0.131)	+0.009 (+0.46, 0.366)	+0.012 (+0.49, 0.333)	+0.010 (+0.80, 0.147)
	NodeCross – NodeRes	+0.006 (+0.26, 0.588)	+0.009 (+0.69, 0.195)	-0.001 (-0.13, 0.787)	+0.003 (+0.33, 0.503)
	GraphCross – GraphRes	-0.003 (-0.30, 0.545)	+0.002 (+0.08, 0.874)	+0.000 (+0.00, 0.999)	-0.008 (-0.32, 0.513)
	GraphRes – NodeRes	+0.015 (+0.87, 0.125)	-0.003 (-0.27, 0.583)	-0.003 (-0.18, 0.714)	+0.005 (+0.23, 0.637)
	NodeCross – GraphCross	-0.005 (-0.35, 0.476)	+0.010 (+0.50, 0.324)	+0.001 (+0.10, 0.832)	+0.007 (+0.49, 0.334)

Table 14. Residual-routing and parameter-sweep summary on four representative targets (transposed). Numbers are five-fold mean accuracy $\pm$ standard deviation. **Bold:** best per column.

Routing	AIDS+GraphRes+GIN	PROT+GraphRes+SAGE	AIDS+NodeCross+GAT	DD+NodeCross+GAT
Learnable	0.9500 $\pm$ 0.0087	0.7178 $\pm$ 0.0411	0.8970 $\pm$ 0.0491	0.7080 $\pm$ 0.0172
Top-k 0.25	0.9075 $\pm$ 0.0237	0.7241 $\pm$ 0.0407	0.9160 $\pm$ 0.0128	0.7062 $\pm$ 0.0249
Top-k 0.50	0.9330 $\pm$ 0.0251	0.7151 $\pm$ 0.0397	0.8960 $\pm$ 0.0483	0.7029 $\pm$ 0.0160
Top-k 0.75	0.9240 $\pm$ 0.0287	0.7214 $\pm$ 0.0318	0.8960 $\pm$ 0.0489	0.7164 $\pm$ 0.0248
Sparse 0.02	0.9470 $\pm$ 0.0181	0.7160 $\pm$ 0.0248	0.9060 $\pm$ 0.0268	0.7147 $\pm$ 0.0259
Sparse 0.05	0.9250 $\pm$ 0.0152	0.7259 $\pm$ 0.0475	0.9180 $\pm$ 0.0162	0.7181 $\pm$ 0.0196
Sparse 0.10	0.8315 $\pm$ 0.0549	0.7088 $\pm$ 0.0521	0.9025 $\pm$ 0.0139	0.7105 $\pm$ 0.0243
Best	Learnable	Sparse 0.05	Sparse 0.05	Sparse 0.05

5 CONCLUSIONS

This paper studies CR-GNN as a controlled biomolecular graph-classification framework and finds that the design space reveals systematic trade-offs among peak accuracy, cross-dataset consistency, and representational continuity. The main contribution is not a universally dominant architecture, but a reproducible comparison framework that makes these trade-offs quantitatively visible and provides practical guidance for architecture selection.

The key empirical findings are as follows. First, GraphResGNN achieves the highest peak accuracy most frequently (50% of 24 dataset–operator combinations) but with the weakest representational continuity at depth (CKA 0.942 at $L = 8$) and the greatest rank volatility (std 1.46). Second, NodeCrossGNN is the second-strongest architecture (25% of combinations), dominating DD across three of four operators, and is particularly effective on large sparse graphs with attention-based operators. Third, GraphCrossGNN is the most rank-consistent design (std 0.75), never falling below fourth place, making it the safest architectural choice when dataset characteristics are uncertain. Fourth, the depth-scaling analysis documents a “CKA–performance paradox”: the architecture with the highest representational continuity (NodeResGNN, CKA 0.968) achieves the fewest wins among reuse variants, while the architecture with the lowest continuity (GraphResGNN) achieves the most. This challenges the assumption that stronger representational stability necessarily produces better performance.

The routing analysis provides mechanism-level evidence that sparse routing combined with cross-residual architectures forms a particularly effective configuration, with sparse mode (threshold 0.05) optimal on three of four representative targets. This synergy—sparsity suppressing weak channels while cross-branch exchange preserves complementary information—represents a practical design principle for

Table 15. Family-wise routing winners.

Target	Best learnable	Best top-k	Best sparse	Overall winner
AIDS + GraphRes + GIN	0.9500	0.9330	0.9470	Learnable
PROTEINS + GraphRes + SAGE	0.7178	0.7241	0.7259	Sparse
AIDS + NodeCross + GAT	0.8970	0.9160	0.9180	Sparse
DD + NodeCross + GAT	0.7080	0.7164	0.7181	Sparse

future GNN architecture development.

The scope of these conclusions should remain disciplined. The strongest claims are supported by the stratified five-fold benchmark, the targeted five-fold ablations, and the depth-scaling analysis ($L = 1-8$) on PROTEINS under GCNConv. The statistical evidence (Tables 12 and 13) confirms that architecture-level effects are modest relative to fold-level variation, consistent with the paper’s mechanism-oriented rather than leaderboard-oriented framing. The study addresses controlled graph-classification behavior on medium-scale biomolecular benchmarks and does not by itself establish large-scale scalability or direct biological interpretability.

In practical terms, this work provides the following guidance: choose GraphResGNN with learnable dense routing when peak accuracy on well-characterized datasets is the priority; choose GraphCrossGNN with sparse routing when consistent, predictable performance across diverse datasets is required; choose NodeCrossGNN with GATConv when working with large, sparse graphs where dual-branch complementary pathways are valuable. These results position CR-GNN as a practical design space and provide a methodological basis for future work on richer biological inputs, larger benchmarks, and graph-learning settings where the trade-offs documented here can be further tested and refined.

6 ACKNOWLEDGMENTS

The authors thank the open-source graph learning community and the maintainers of the benchmark resources used in this study. They also thank the reviewers and editors for their time and feedback.

REFERENCES

- Uri Alon and Eran Yahav. On the bottleneck of graph neural networks. In *International Conference on Learning Representations (ICLR)*, 2021.
- Xavier Bresson and Thomas Laurent. Residual gated graph convnets. *arXiv preprint arXiv:1711.07553*, 2017.
- Tianle Cai, Jiacheng Luo, Sheng Wang, Kai Zhang, Yu Tian, Liwei Yang, Zekun Li, and Kilian Weinberger. Graphnorm: A principled approach to accelerating graph neural network training. In *International Conference on Machine Learning*, pages 1277–1287, 2021.
- Ming Chen, Zhewei Wei, Zengfeng Huang, Bolin Ding, and Yaliang Li. Simple and deep graph convolutional networks. In *International Conference on Machine Learning (ICML)*, pages 1725–1735, 2020a.
- Yuning Chen, Xihao Wei, Pan Xu, Xinwen Wang, Ping Luo, Zhiyuan Li, and Jian Tang. Revisiting over-smoothing in deep gcns. *arXiv preprint arXiv:2003.13663*, 2020b.
- Matthias Fey and Jan Eric Lenssen. Fast graph representation learning with pytorch geometric. *arXiv preprint arXiv:1903.02428*, 2019.
- Hongyang Gao and Shuiwang Ji. Graph u-net. *arXiv preprint arXiv:1905.05178*, 2019.
- Justin Gilmer, Samuel S Schoenholz, Patrick F Riley, Oriol Vinyals, and George E Dahl. Neural message passing for quantum chemistry. In *International Conference on Machine Learning*, pages 1263–1272, 2017.
- William L Hamilton, Rex Ying, and Jure Leskovec. Inductive representation learning on large graphs. In *Advances in neural information processing systems*, pages 1024–1034, 2017.

522 Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition.
523 In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778,
524 2016.

525 Thomas N Kipf and Max Welling. Semi-supervised classification with graph convolutional networks.
526 *arXiv preprint arXiv:1609.02907*, 2016.

527 Johannes Klicpera, Alan Bojchevski, and Stephan Günnemann. Combining label propagation and simple
528 models out-performs graph neural networks. *arXiv preprint arXiv:1810.05997*, 2018.

529 Guohao Li, Elias Muller, Abbas Thabet, and Bernard Ghanem. Deeper insights into graph convolutional
530 networks for semi-supervised learning. *arXiv preprint arXiv:1809.02584*, 2018.

531 Christopher Morris, Roger Wattenhofer, and Kristian Kersting. Tudataset: A collection of benchmark
532 datasets for learning with graphs. *arXiv preprint arXiv:2007.08663*, 2020.

533 Kenta Oono and Taiji Suzuki. Simple and deep graph convolutional networks. *arXiv preprint*
534 *arXiv:2006.13604*, 2020.

535 Yu Rong, Wenbing Huang, Tingyang Xu, and Junzhou Huang. Dropedge: Towards deep graph convolu-
536 tional networks on node classification. *arXiv preprint arXiv:1907.10903*, 2019.

537 Jake Topping, Francesco Di Giovanni, Benjamin Chamberlain, Michael Bronstein, Douwe Vendrig, and
538 Pietro Lio. Understanding over-squashing and bottlenecks on graphs via curvature. *arXiv preprint*
539 *arXiv:2111.14522*, 2021.

540 Zonghan Wu, Shirui Pan, Fengwen Chen, Guoding Long, Cheng Zhang, and Sheng Yu Zhang. A
541 comprehensive survey on graph neural networks. *IEEE Transactions on Neural Networks and Learning*
542 *Systems*, 32(1):4–24, 2021.

543 Keyulu Xu, Chengtao Li, Yonglong Tian, Xiao Du, Jure Leskovec, and Stefanie Jegelka. Representation
544 learning on graphs with jumping knowledge networks. In *International Conference on Machine*
545 *Learning*, pages 5449–5458, 2018.

546 Rex Ying, Jiaxuan You, Christopher Morris, Xiang Ren, William L Hamilton, and Jure Leskovec. Hierar-
547 chical graph representation learning with differentiable pooling. In *Advances in neural information*
548 *processing systems*, pages 4800–4810, 2018.

549 Lingxiao Zhao and Leman Akoglu. Pairnorm: Tackling over-smoothing in gns. *arXiv preprint*
550 *arXiv:2009.10698*, 2020.

551 Jie Zhou, Ganqu Cui, Zhengyan Zhang, Cheng Yang, Zhaocheng Liu, Zhongyuan Wang, Peng Li, Ming
552 Sun, Yu Shi, et al. Graph neural networks: A review of methods and applications. *AI Open*, 1:57–81,
553 2020.